

Agentic Commerce Architecture

High-level design for a system in which a software agent can spend a human's money — bounded by rules the agent cannot reach, and recorded so that every rupee has an explanation.

The policy engine is the product. The agent only proposes.

A language model decides what to attempt; deterministic Python decides what is permitted.

Track 01 — AI Growth & Agentic Commerce

Stack FastAPI · Next.js · Agno · NVIDIA NIM · Razorpay (test mode)

Verification 303 tests · 13 policy rules · 36 audit event types

Architecture

High-level design for an agentic-commerce system in which a software agent can spend a human's money — bounded by rules the agent cannot reach, and recorded so that every rupee has an explanation.

Audience: an engineer who wants to understand how this is put together without reading the code.


The thesis

The policy engine is the product. The agent only proposes.

A language model decides *what to attempt*; deterministic Python decides *what is permitted*. The model never executes anything. It emits a proposed tool call, which the policy engine evaluates against the catalog, the cart, the session budget and the acting credential's own spend limit — and returns `ALLOW`, `REQUIRE_CONFIRMATION` or `DENY`. Nothing about the model's confidence, phrasing or reasoning enters that decision. If the model is jailbroken, hallucinating, or replaced tomorrow by a better one, the bounds on money are unchanged, because they were never expressed to the model in the first place.

Everything below is in service of that split.

1. System context



What crosses each boundary

BOUNDARY	DIRECTION	CARRIES	TRUST
Human to platform	in	Google-issued identity, then our own JWT	Authenticated
Agent to platform	in	X-Agent-Key , hashed at rest	Authorized, never authenticated
Platform to Razorpay	out	order creation, signature verification	External authority on money
Platform to NVIDIA	out	conversation + tool schemas	Untrusted output
Platform to agent	out	catalog, prices, policy decisions with reasons	Public / aggregate only

The asymmetry in row two is the point. **A human authenticates — they prove who they are. An agent is authorized — a human grants it a bounded capability.** An agent cannot create another agent, cannot read its own scopes or spend limit, and cannot reach any buyer-only endpoint.

The NVIDIA row is the other load-bearing one: **model output is data, never instruction.** Tool results are never spliced into the system prompt, and every proposed action is re-checked by the policy engine regardless of what the model produced.

2. Layer view

The main structural diagram. Read top to bottom; each layer may only call the one below it.




Figure 2a — the request path: client to governance. Green is deterministic and load-bearing; orange is the only layer where a model is involved.




Figure 2b — domain, persistence, and the cross-cutting audit trail.


Green is deterministic and load-bearing. Orange is the only layer where a model is involved.
Nothing in green reads model output as instruction.

Layer responsibilities

#	LAYER	RESPONSIBLE FOR	MUST NEVER	AT SCALE
1	Client	Rendering, collecting intent	Hold authority; every check is re-run server-side	CDN; unchanged otherwise
2	Edge/auth	Deciding <i>who</i> is calling; failing closed	Let an unmarked route through	Move JWT verify to a gateway; add rate limiting
3	Application	Use-case orchestration	Routers never touch the DB — SQL lives only in repositories	Split read/write paths
4	Agent	Turning language into <i>proposals</i>	Execute anything itself; trust its own output	Queue turns; stream partials
5	Governance	Deciding what is permitted	Consult a model, or be bypassable	Unchanged — this is pure functions
6	Domain	Cart, order, payment, campaign, demand rules	Skip the state machine or write payments directly	Extract payments first
7	Data	Persistence	Hold business logic	Postgres; the code already supports it
8	Audit	Recording every decision with its reason	Offer any update or delete	Async writer; separate store

The layering rule that does the most work is layer 3: **routers never touch the database**. A router parses HTTP and delegates. That is why the policy engine cannot be accidentally bypassed by a new endpoint — there is no path from HTTP to the database that does not pass through a service.

3. Request lifecycle — one purchase, end to end



Two things this diagram is drawn to make obvious:

1. **The model's output is a proposal that passes through the green band before anything happens.**
2. **The cart is re-validated at payment time**, not trusted from when items were added — prices, stock and limits are all re-checked against the state of the world *now*.

4. Data model




Figure 4a — identity and commerce.





Figure 4b — agent sessions and merchant growth.

Grouped by concern:

GROUP	TABLES	NOTE
Identity	users , agent_credentials	An agent is never a User row. Only key_hash is stored
Commerce	products , carts , cart_items , orders , payments	Money is integer pause, everywhere
Agent	agent_sessions , agent_messages	Full transcript; a session belongs to one credential
Growth	demand_signals , merchant_notifications , product_views , campaign_*	demand_signals has no user_id column at all
Audit	audit_events	Append-only. 36 event types

Two deliberate absences. demand_signals carries no buyer identity — "distinct buyers" is approximated by distinct session, so a merchant-facing aggregate cannot be de-anonymised, because the identity is not in the table to select. And audit_events has no update or delete path in the repository at all.


5. Order state machine



FAILED → PAID was added after order #23 was recorded **FAILED** while Razorpay held a real captured ₹275 payment — Checkout reported failure and success out of order within one session. The reverse remains forbidden.

6. External AI buyer path

The path that involves no human UI at all.



`/.well-known/catalog.json` is a discovery document: capabilities, endpoint URLs, and a plain-English `how_to_transact`. It borrows the `/.well-known/` convention from UCP and A2A, and ACP's field names where they apply — stated precisely in the document itself rather than claiming conformance.

7. Four guarantees that hold by construction

Not by convention, review, or discipline — by there being no code path to the alternative.

GUARANTEE	MECHANISM	WHERE
The audit log cannot be rewritten	The repository exposes <code>create</code> plus three reads. There is no update or delete method to call	<code>app/audit/repository.py</code>
An agent can never complete a payment alone	<code>PaymentAuthorizationRule</code> has no branch that returns <code>ALLOW</code> — only <code>REQUIRE_CONFIRMATION</code> or <code>DENY</code>	<code>app/policy/rules.py</code>
Chaos and demo login cannot exist in production	Gated on <code>settings.app_env == "development"</code> in code. No env var, header or request body can enable them elsewhere	<code>app/testing/chaos.py</code> , <code>app/testing/demo_login.py</code>
Campaign control groups cannot be gamed	The control/treatment split happens <i>before</i> policy evaluation and is written as a <code>control_group_split</code> audit event	<code>app/campaigns/service.py</code>

The first two are the ones worth arguing about, so they are stated as testable properties rather than intentions: *there is no delete method*, and *there is no ALLOW path*. Both are verifiable by grep in under a minute.

8. Policy engine internals

Every rule is a class with one method:

```
class Rule(ABC):
    def evaluate(self, action: ProposedCartState) -> RuleResult | None: ...
```

`None` means "this rule has no opinion". A `RuleResult` carries a decision, the rule's own name, and a **human-readable reason built from the actual numbers** — that reason is what reaches the buyer and the audit trail.

Resolution is simple and total: `DENY` beats `REQUIRE_CONFIRMATION` beats `ALLOW`. Within a decision, registration order breaks ties, so ordering encodes priority — credential checks before cart arithmetic.

The 13 registered rules:

GROUP	RULES
Credential	<code>RevokedCredentialRule</code> , <code>AgentScopeRule</code> , <code>AgentSpendLimitRule</code>
Catalog	<code>UnknownSkuRule</code> , <code>OutOfStockRule</code> , <code>StockRule</code>
Limits	<code>PerItemPriceRule</code> , <code>QuantityRule</code> , <code>SpendCapRule</code>
Offers	<code>UpsellPolicyRule</code>
Payment	<code>ConfirmationThresholdRule</code> , <code>DuplicatePaymentRule</code> , <code>PaymentAuthorizationRule</code>

Adding a rule is three steps: implement `evaluate`, register it in `default_policy_engine()` at the right priority, add a test. No other file changes. Nothing needs to know the rule exists.

Why deterministic Python and never prompt instructions. A prompt is a request; a rule is a constraint. "Never spend more than ₹500" in a system prompt is advice the model may ignore, misread under adversarial input, or silently drop when the prompt is edited. As Python it is a branch — and it is re-

evaluated at payment time against the cart as it exists then, not as it existed when the model was persuaded. This also makes the bound *auditable*: a reviewer can read thirteen small classes and know the complete set of things that can stop a purchase.

`PaymentAuthorizationRule` is worth reading closely: rather than duplicating thresholds, it replays every cart line through the item-level rules as if being freshly added, and surfaces the underlying rule's name on failure — so the audit line says `SpendCapRule`, not a useless wrapper name.

9. Where AI is used, and where it deliberately isn't

This is a design position, not an omission.

AI is used for exactly three things

USE	WHY A MODEL EARNS ITS PLACE
Intent to tool calls	Mapping "5kg atta under ₹400" onto a tool call with arguments is genuinely open-ended language understanding
Demand-signal extraction	Pulling a category and constraints out of free text, where the phrasing space is unbounded
Campaign copy + rationale	Writing an offer message is a language task with no correct answer

AI is deliberately *not* used for

TASK	WHAT IS USED INSTEAD	WHY
Every money decision	The policy engine	A bound that a model can be talked out of is not a bound
Customer segmentation	Deterministic SQL + Python	"3+ orders" is a definition, not a judgement. An LLM would add cost, latency and non-reproducibility for nothing
Offer sizing / discount caps	<code>DiscountCapRule</code> , <code>MarginFloorRule</code>	Margin protection must be exact and provable
Demand thresholds	<code>max(5, active_buyers x 0.20)</code>	A threshold is arithmetic
Notification aggregation	Grouped counts	Deterministic and idempotent — safe to recompute on every read
Product search over 50 SKUs	SQL <code>LIKE</code> + tags	A vector DB for 50 products is infrastructure cost with no accuracy gain
Order state transitions	A pure state machine	Correctness here is a table, not an inference

One campaign path records this explicitly in its own audit trail: `model_used = "none (deterministic browse-abandonment offer)"`.

The rule of thumb applied throughout: if the answer is a definition, arithmetic, or a table, it is code. If the answer is open-ended language, it is a model — and its output is then treated as an untrusted

proposal.

10. Key design decisions

DECISION	ALTERNATIVE REJECTED	REASON
Policy engine as code	Rules in the system prompt	A prompt is advice; a branch is a constraint. Also auditable and testable
<code>Model.invoke()</code>	Agno's <code>Agent.run()</code>	<code>Agent.run()</code> auto-executes tool calls , which would bypass the policy gate entirely. <code>invoke()</code> returns one proposal and executes nothing — the harness stays in control
Integer paise	Float rupees	Float money accumulates rounding error. Every amount in the system is an integer
No vector DB	pgvector / Chroma	50 products. <code>LIKE</code> + tags is faster, exact, and has no index to drift
Agents authorized, not authenticated	Give agents user accounts	An agent is a <i>capability a human grants</i> , revocable independently, with its own scopes and cap. A user account would make it a peer of its owner
Append-only audit by construction	Soft-delete flag	A flag is a convention someone can bypass. No delete method is a property
SQLite by default	Postgres required	One command from clone to running. The code is engine-agnostic; Postgres is a compose profile
Idempotency includes <code>cart_id</code>	Content hash only	Without it, re-ordering an identical basket collided with the previous purchase and was refused forever

11. Non-functional requirements

Requirement	How it is met	Honest status
Idempotency	Content + cart hash with a DB <code>UNIQUE</code> constraint; insert-then-catch, never check-then-insert	Met
Observability	36 audit event types; structured JSON logs with <code>request_id</code> threaded through; session replay	Met
Reliability	Retry, fallback model, circuit breaker; 7 injectable faults	Met
Security	Default-deny routing; HMAC verified server-side; keys hashed at rest; scoped credentials	Met for the modelled threats
Data integrity	Pure state machine; integer money; append-only audit	Met
Latency	Catalog reads are ms. Agent turns are 5-120s , bounded by the model provider	Not met, and not solvable here — it is provider latency, visible in the UI as a pending state
Availability	Single node, single process	Not a goal for this build

12. Failure modes


FAILURE	RESPONSE	RECORDED AS
Model unavailable	Retry, fallback model, circuit breaker fails fast for a cooldown	<code>model_call_failed</code>
Model emits malformed tool call	Rejected before execution	<code>malformed_tool_call</code>
Model invents a SKU	<code>UnknownSkuRule</code> denies	<code>policy_decision</code>
Model loops	Hard cap at 8 iterations	<code>iteration_limit_hit</code>
Injection text in catalog data	Flagged; policy re-checks the action regardless	<code>injection_detected</code>
Razorpay unreachable	Order marked <code>FAILED</code> , retryable	<code>razorpay_order_failed</code>
Tampered payment signature	Refused; order not marked paid	<code>signature_rejected</code>
Duplicate payment attempt	Existing order returned, no second charge	<code>duplicate_payment_prevented</code>
Repeat failure callback	Idempotent; no 500	<code>payment_failed</code>
Late verified success after failure	Recovered to <code>PAID</code> with an explicit event	<code>payment_recovered_after_failure</code>
Unhandled server error	500 with CORS headers , so the browser can read it	<code>unhandled_exception</code> + traceback

13. Scaling analysis — what breaks first

In order:

1. **SQLite, under concurrent writes.** Single writer. First thing to hit. Mitigation exists today: `DATABASE_URL` already accepts Postgres and a compose profile ships it.
2. **Synchronous audit writes.** Every decision writes a row in the request path. At volume this becomes the dominant cost of a turn. Fix: async writer or an append-only log store — the repository interface is narrow enough that this is a contained change.
3. **No `payment.captured` webhook.** The system learns about payments from a browser callback. If the buyer closes the tab after paying, Razorpay holds a captured payment the system never hears about. **This is the most important architectural gap**, and it is why four orders remain unreconciled. A webhook is the correct answer and is not built.
4. **Single node, in-process background work.** Title generation runs on a thread. Anything more needs a real queue.
5. **Per-request model latency.** 5–120s per agent turn. Unfixable in-process; needs streaming or a job model.

14. Security model




Principal types and their reach:

- **PENDING** — signed in with Google, no role chosen. Reaches onboarding and nothing else.
- **BUYER** — own cart, own orders, own agents, chat.
- **MERCHANT** — catalog, campaigns, notifications, audit. No access to any buyer's cart.
- **AGENT** — a scoped subset, bounded by `spend_limit_paise`. Cannot mint credentials, cannot read its own configuration, cannot reach buyer-only endpoints.

Model output and merchant-authored catalog text are both treated as untrusted input.

15. The supervisor — designed, not built

The natural next layer, described here because its absence is a real limitation rather than an oversight.



The gap it would close: per-credential limits do not compose. Each agent is correctly bounded at ₹500. Three agents are three separate ₹500 bounds — **₹1,500 of account exposure that no single rule sees**, because every rule evaluates one credential's action in isolation. Nothing in the system today enforces a ceiling across an owner's agents.

A supervisor would add an account-level aggregate cap, velocity limits across credentials, and anomaly detection on spending patterns. It is deliberately named rather than quietly omitted: an agentic-commerce system that claims bounded spend should be explicit about the boundary its bounds do not cross.

Reading order for the code

TO UNDERSTAND	READ
The thesis in one file	<code>app/policy/rules.py</code>
How an agent is stopped	<code>app/policy/engine.py</code>
Why a model can't execute anything	<code>app/agent/harness.py</code>
Default-deny	<code>app/auth/routing.py</code>
What is real at Razorpay	<code>app/payments/gateway.py</code> , then PAYMENT-REALITY.md
What broke	Failures.md
What was verified	SYSTEM-AUDIT.md